


Generalized File System Dependencies

Christopher Frost* Mike Mammarella* Eddie Kohler*
Andrew de los Reyes† Shant Hovsepian*
Andrew Matsuoka‡ Lei Zhang†

*UCLA †Google ‡UT Austin

<http://featherstitch.cs.ucla.edu/>



Featherstitch Summary

- A **new architecture** for constructing file systems
- The **generalized dependency** abstraction
 - Simplifies consistency code within file systems
 - Applications can **define consistency requirements** for file systems to enforce

File System Consistency

- Want: don't lose file system data after a crash
- Solution: keep file system consistent after every write
 - Disks do not provide atomic, multi-block writes
- Example: journaling



- Enforce *write-before relationships*

File System Consistency Issues

- Durability features vs. performance
 - Journaling, ACID transactions, WAFL, soft updates
 - Each file system picks one tradeoff
 - Applications get that tradeoff plus sync
- Why no extensible consistency?
 - Difficult to implement
 - Caches complicate write-before relations
 - Correctness is critical

“Personally, it took me about 5 years to thoroughly understand soft updates and I haven't met anyone other than the authors who claimed to understand it well enough to implement it.” – Valerie Henson

FreeBSD and NetBSD have each recently attempted to add journaling to UFS. Each declared failure.

The Problem

Can we develop a simple, general mechanism for implementing *any* consistency model?

Yes! With the *patch* abstraction in Featherstitch:

- File systems specify low-level write-before requirements
- The buffer cache commits disk changes, obeying their order requirements

Featherstitch Contributions

- The *patch* and *patchgroup* abstractions
 - Write-before relations become explicit and file system agnostic
- Featherstitch
 - *Replaces* Linux's file system and buffer cache layer
 - ext2, UFS implementations
 - Journaling, WAFL, and soft updates, implemented using just patch arrangements
- Patch optimizations make patches practical

Patches

Problem

Patches for file systems

Patches for applications

Patch optimizations

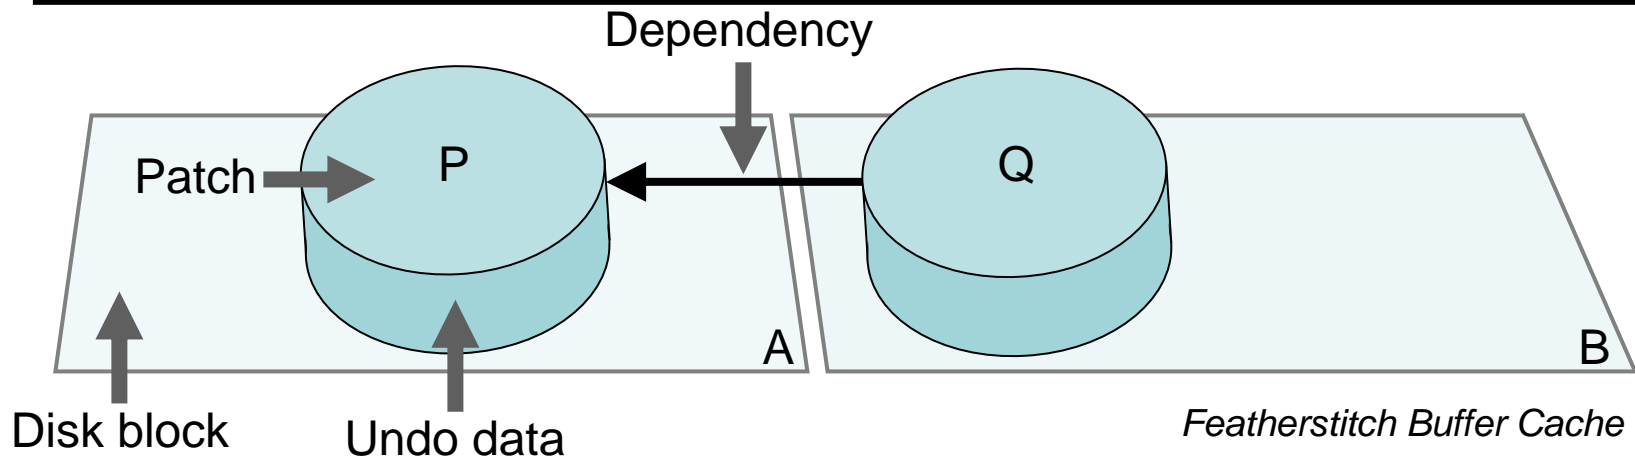
Evaluation

Patch Model

A patch represents:

- a disk data change
- any dependencies on other disk data changes

```
patch_create(block* block, int offset, int length, char* data, patch* dep)
```





Benefits:

- separate write-before specification and enforcement
- explicit write-before relationships

Base Consistency Models


- Fast
 - Asynchronous
- Consistent
 - Soft updates
 - Journaling
- Extended
 - WAFL
 - Consistency in file system images
- All implemented in Featherstitch

Patch Example: Asynchronous rename()



File lost.


Patch Example: rename() With Soft Updates




A valid block writeout:

time →


Patch Example: rename() With Soft Updates




Block level cycle:



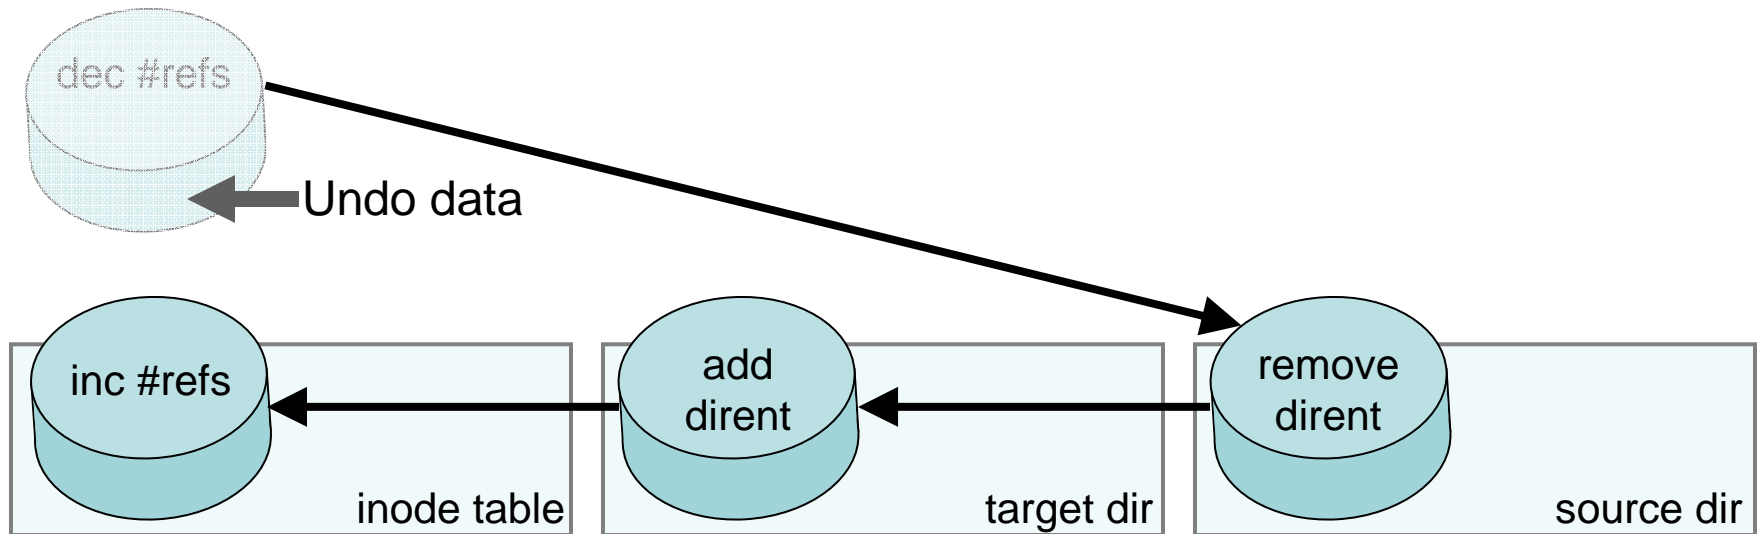
Patch Example: rename() With Soft Updates



Not a *patch level* cycle:



Patch Example: rename() With Soft Updates

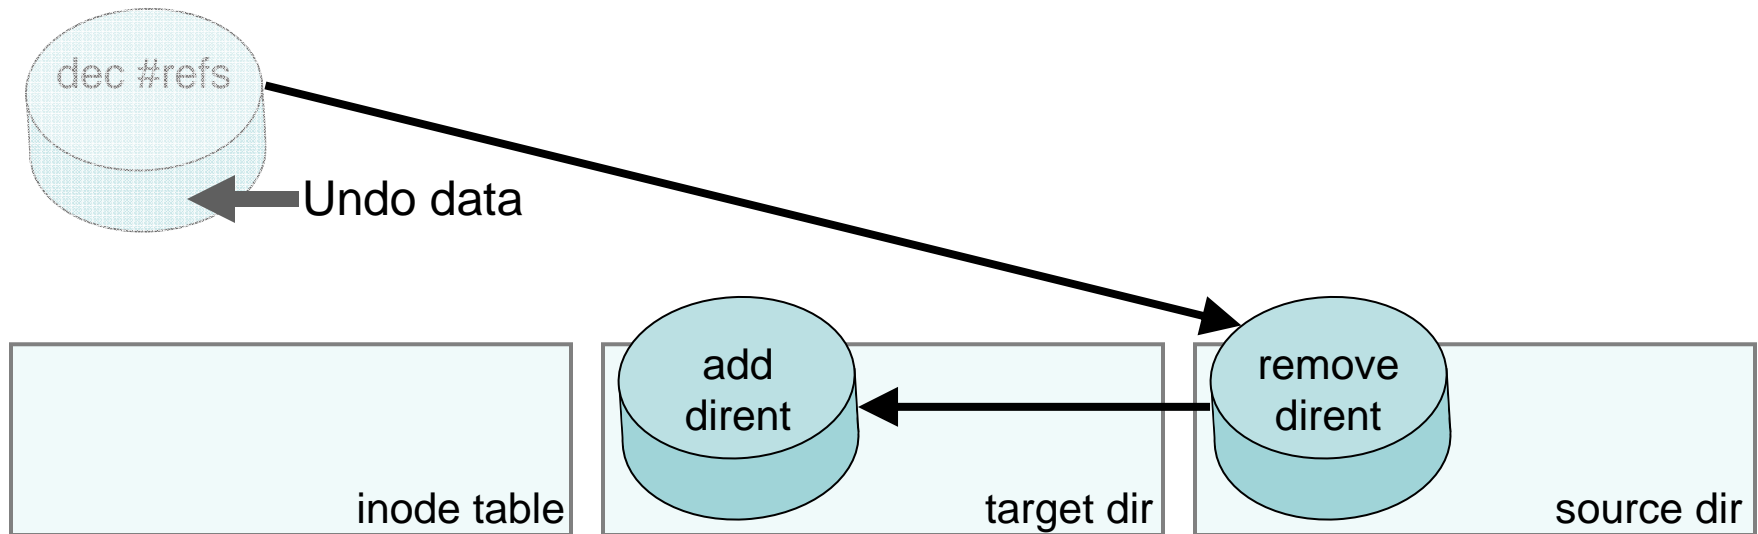


A valid block writeout:



time →

Patch Example: rename() With Soft Updates




A valid block writeout:




time →

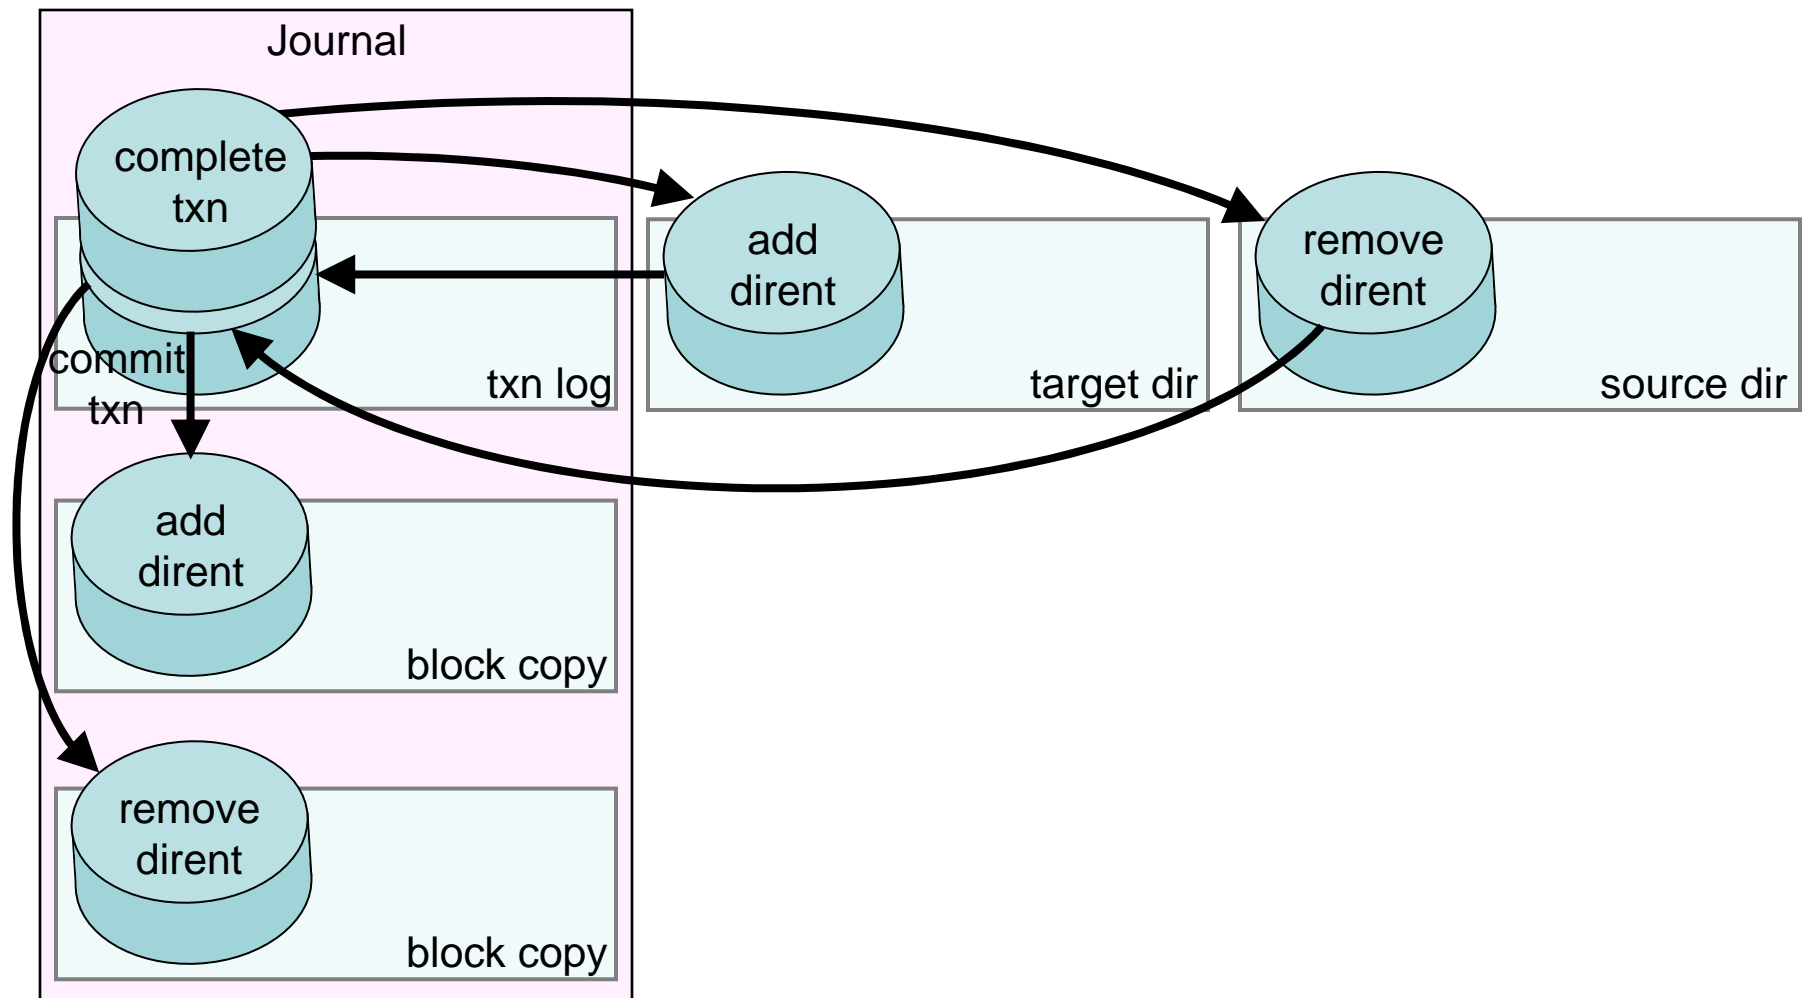
Patch Example: rename() With Soft Updates



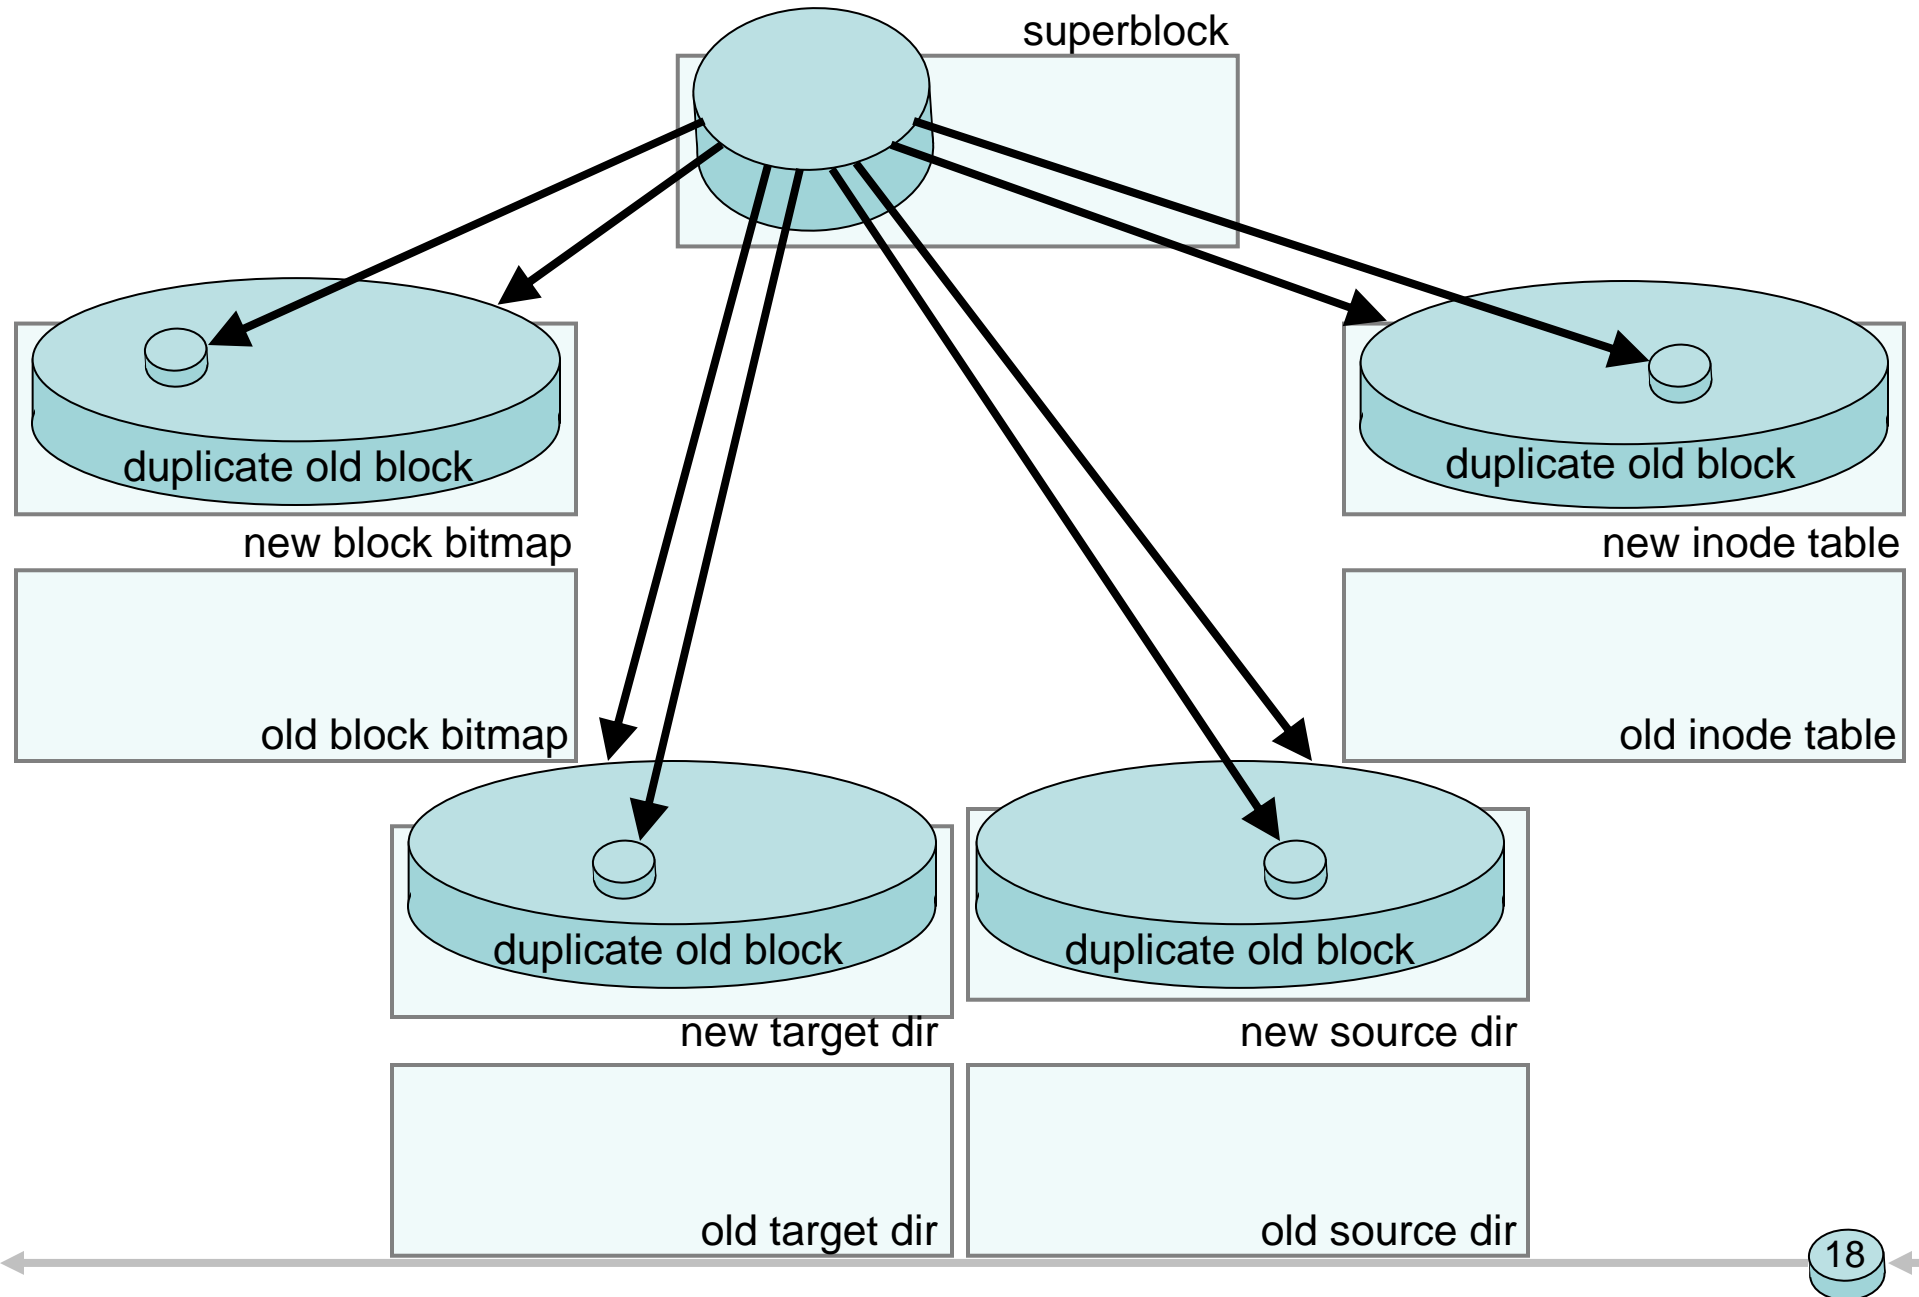
A valid block writeout:



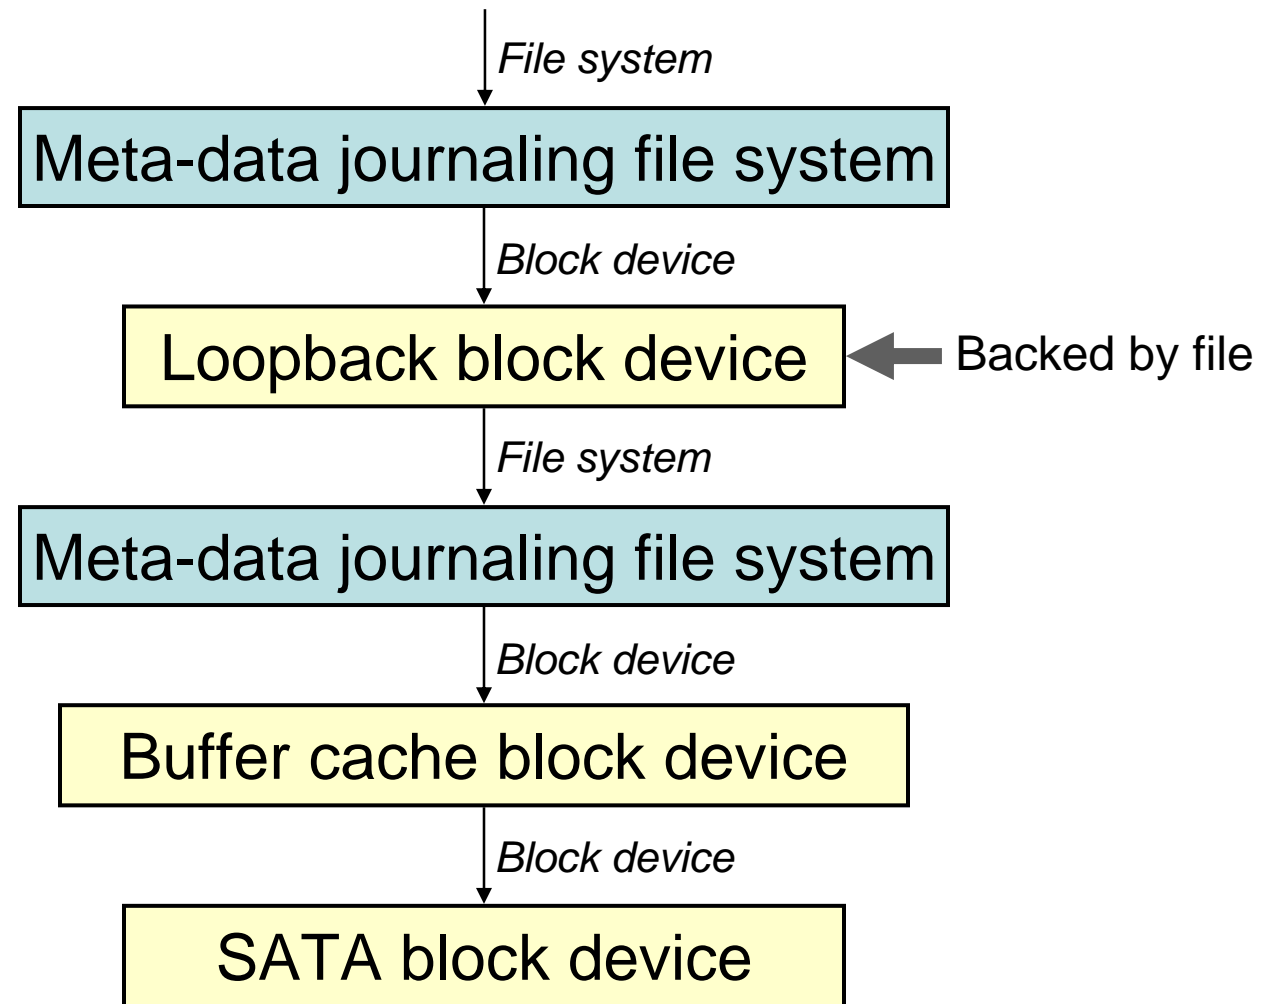
Patch Example: rename() With Journaling



Patch Example: rename() With WAFL



Patch Example: Loopback Block Device



Meta-data journaling file system obeys file data requirements

Patchgroups

Problem

Patches for file systems

Patches for applications

Patch optimizations


Evaluation

Application Consistency

- Application-defined consistency requirements
 - Databases, Email, Version control
- Common techniques:
 - Tell buffer cache to write to disk immediately (fsync et al)
 - Depend on underlying file system (e.g., ordered journaling)

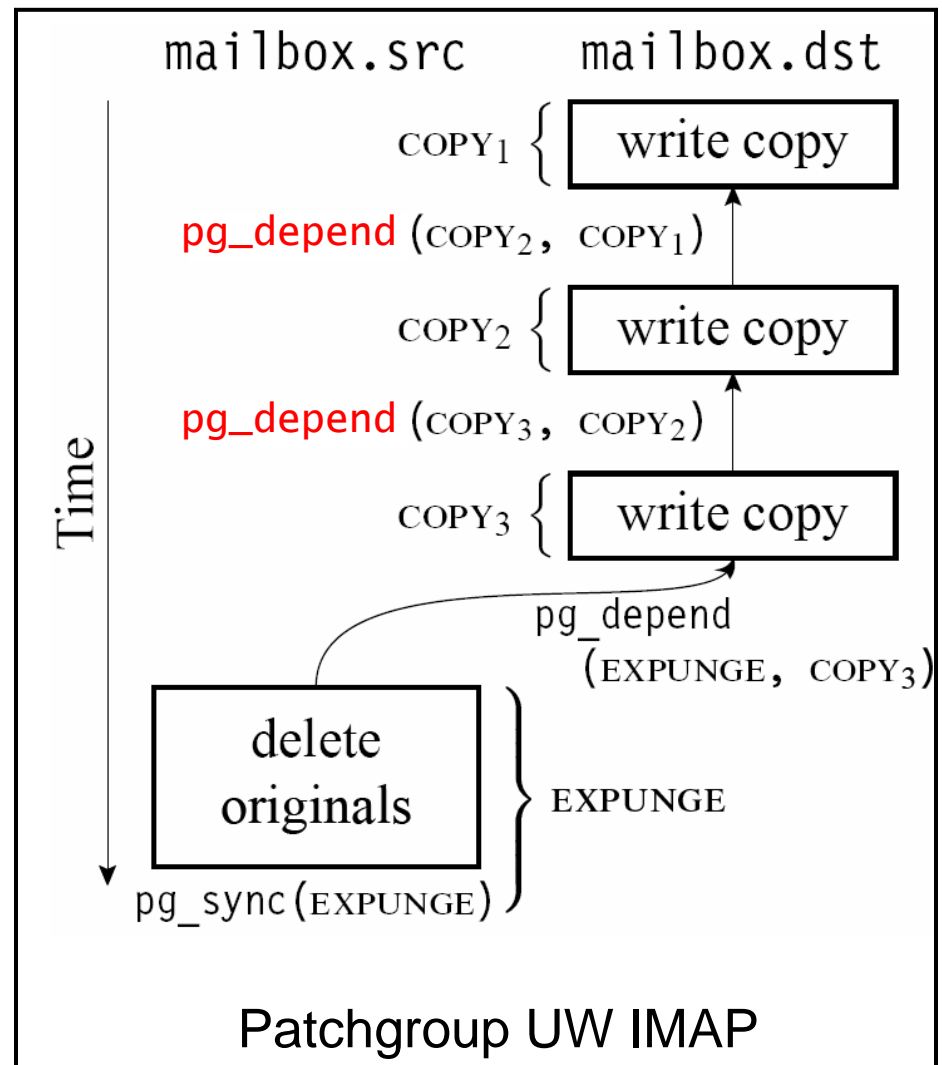
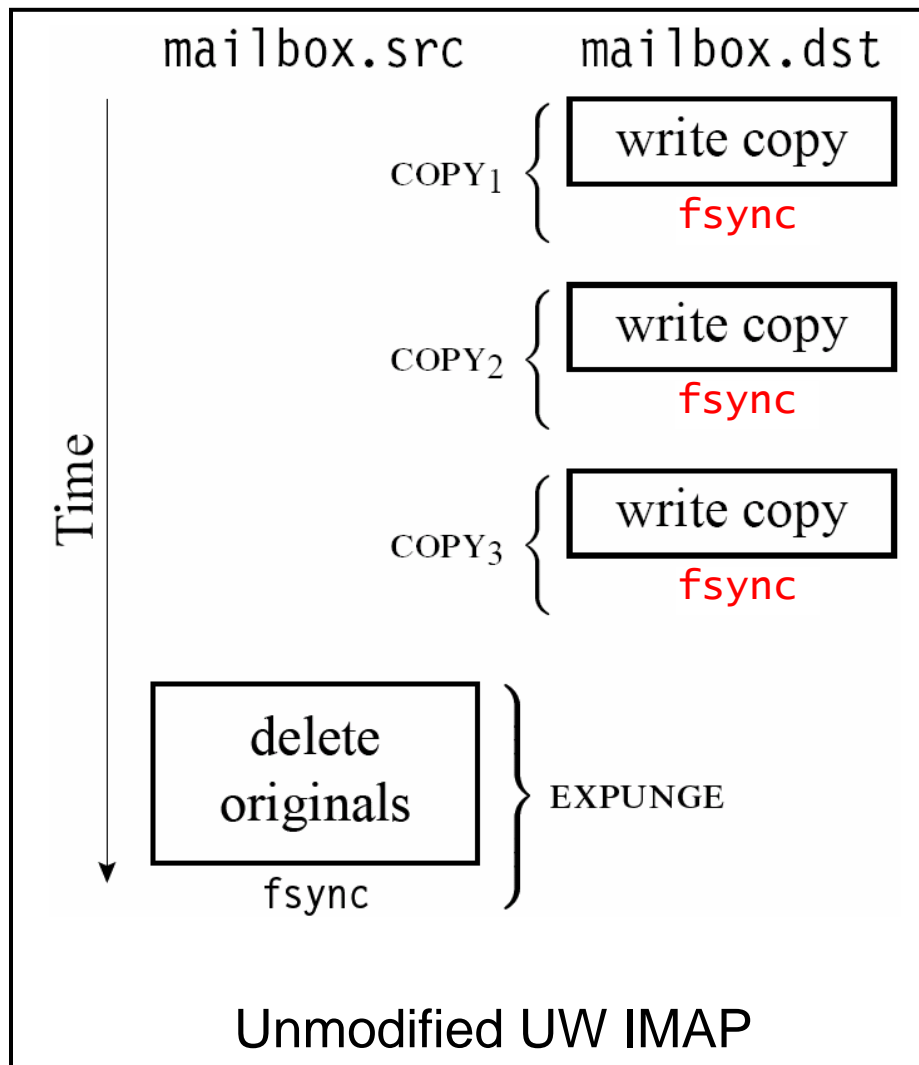
Patchgroups

- Extend patches to applications: *patchgroups*
 - Specify write-before requirements among system calls



- Adapted gzip, Subversion client, and UW IMAP server

Patchgroups for UW IMAP



Patch Optimizations

Problem

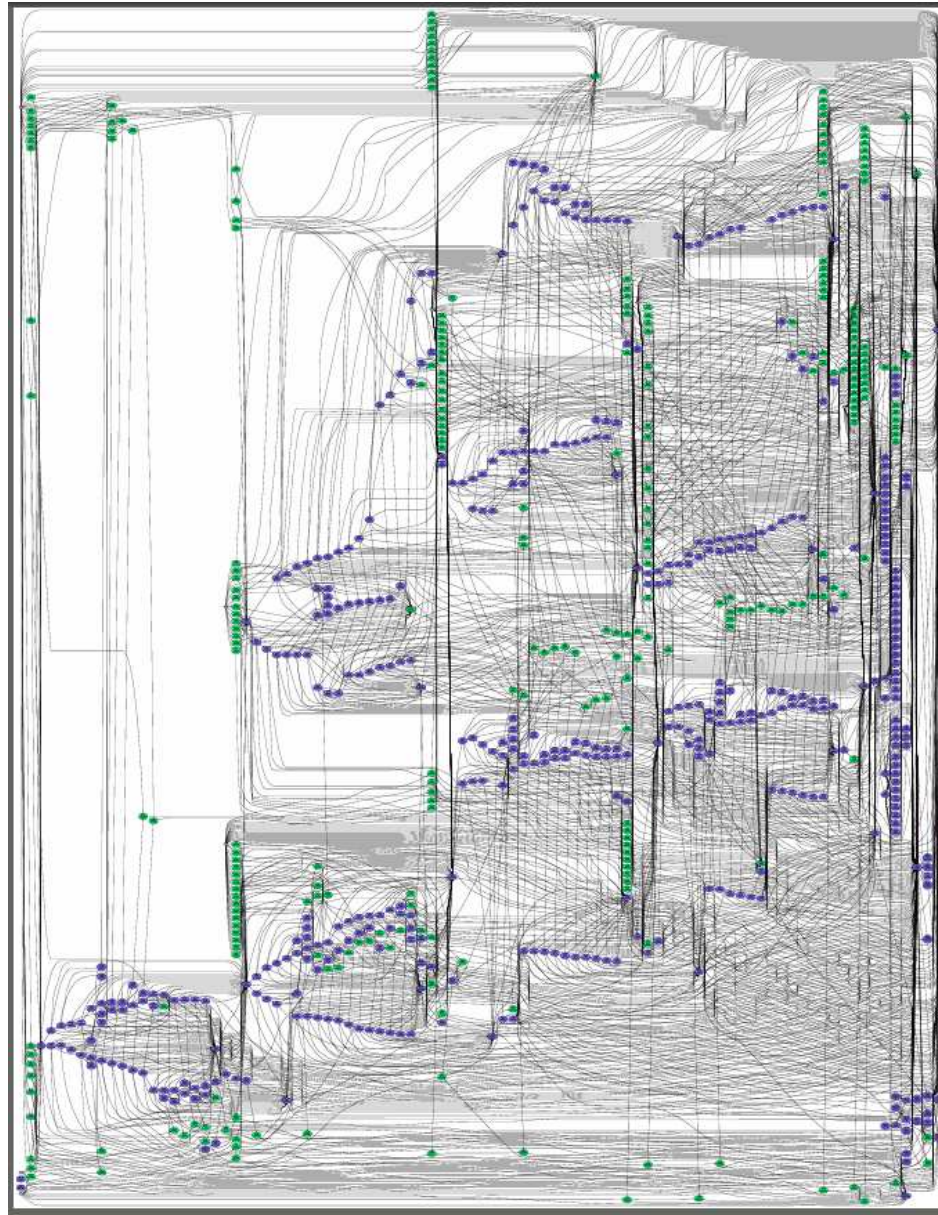
Patches for file systems

Patches for applications


Patch optimizations

Evaluation

Patch Optimizations



Patch Optimizations





- In our initial implementation:
 - Patch manipulation time was the system bottleneck
 - Patches consumed more memory than the buffer cache
- File system agnostic patch optimizations to reduce:
 - Undo memory usage
 - Number of patches and dependencies
- Optimized Featherstitch is not much slower than Linux ext3

Optimizing Undo Data

- Primary memory overhead: unused (!) undo data
- Optimize away unused undo data allocations?
 - Can't detect “unused” until it's too late
- Restrict the patch API to reason about the future?



Optimizing Undo Data

Theorem: A patch that must be reverted to make progress must *induce a block-level cycle*.




Hard Patches

- Detect block-level cycle inducers when allocating?
 - Restrict the patch API: supply all dependencies at patch creation*
- Now, any patch that will need to be reverted must induce a block-level cycle at creation time
- We call a patch with undo data omitted a **hard patch**. A **soft patch** has its undo data.




Patch Merging

- Hard patch merging



- Overlap patch merging



Evaluation

Problem

Patches for file systems


Patches for applications

Patch optimizations

Evaluation

Efficient Disk Write Ordering

- Featherstitch needs to efficiently:
 - Detect when a write becomes durable
 - Ensure disk caches safely reorder writes
- SCSI TCQ or modern SATA NCQ
+ FUA requests or WT drive cache
- Evaluation uses disk cache safely for both
Featherstitch and Linux



Evaluation

- Measure patch optimization effectiveness
- Compare performance with Linux ext2/ext3
- Assess consistency correctness
- Compare UW IMAP performance

Evaluation: Patch Optimizations

PostMark

Optimization	# Patches	Undo data	System time
None	4.6 M	3.2 GB	23.6 sec
Hard patches	2.5 M	1.6 GB	18.6 sec
Overlap merging	550 k	1.6 GB	12.9 sec
Both	675 k	0.1 MB	11.0 sec

Evaluation: Patch Optimizations

PostMark

Optimization	# Patches	Undo data	System time
None	4.6 M	3.2 GB	23.6 sec
Hard patches	2.5 M	1.6 GB	18.6 sec
Overlap merging	550 k	1.6 GB	12.9 sec
Both	675 k	0.1 MB	11.0 sec


Evaluation: Linux Comparison

Fstitch total time

Fstitch system time

Linux total time

Linux system time



- Faster than ext2/ext3 on other benchmarks
 - Block allocation strategy differences dwarf overhead

Evaluation: Consistency Correctness

- Are consistency implementations correct?
- Crash the operating system at random
- Soft updates:
 - Warning: High inode reference counts (**expected**)
- Journaling:
 - Consistent (**expected**)
- Asynchronous:
 - Errors: References to deleted inodes, and others (**expected**)

Evaluation: Patchgroups

- Patchgroup-enabled vs. unmodified UW IMAP server
benchmark: move 1,000 messages
- Reduces runtime by 50% for SU, 97% for journaling

Related Work

- Soft updates [Ganger '00]
- Consistency research
 - WAFL [Hitz '94]
 - ACID transactions [Gal '05, Liskov '04, Wright '06]
- Echo and CAPFS distributed file systems [Mann '94, Vilayannur '05]
- Asynchronous write graphs [Burnett '06]
- xsyncfs [Nightingale '05]

Conclusions

- Patches provide new write-before abstraction
- Patches simplify the implementation of consistency models like journaling, WAFL, soft updates
- Applications can precisely and explicitly specify consistency requirements using patchgroups
- Thanks to optimizations, patch performance is competitive with ad hoc consistency implementations

Featherstitch source:

<http://featherstitch.cs.ucla.edu/>

Thanks to the NSF, Microsoft, and Intel.

